

How do we know if a procedure “works” properly?

- Does a test give the right p-values?
- Do confidence intervals have the right coverage rates?

A typical stat article claims some new procedure is better than an old one. Arguments:

- likelihood based – often using MLE’s are asymptotically normal. BUT how big does n need to be to use this argument?
- simulation based – create fake data based on an assumed parameter value.
 - Show that the procedure estimates the parameter and gives good coverage. (Repeat process many times. Examine the proportion of parameters covered by their estimating intervals.)
 - May use different n to see how it works for small sample sizes as well as large ones.

8.1 SLR

Assume we have model $y_i = \alpha + \beta x_i + \epsilon_i$

Create fake data, check coverage rates of CI’s for β .

- 1 Arbitrary “true” values:

$$\alpha = 1.4, \quad \beta = 2.3, \quad \sigma = 0.9, \quad \mathbf{x} = (1 \ 2 \ 3 \ 4 \ 5)^T$$

```
alpha <- 1.4
beta <- 2.3
sigma <- 0.9
x <- 1:5
n <- length(x)
```

- 2 Generate data.

```
y <- alpha + beta * x + rnorm(n, 0, sigma)
```

- 3 Forget the parameters, and estimate them.

	Estimate	Std. Error	t value
(Intercept)	-0.63	0.29	-2.17
x	2.76	0.09	31.63

Table: $n = 5$ rank = 2 resid sd = 0.276 R-Squared = 0.997

A powerful tool, needs some skill and practice for application
Not a “standard” procedure (unlike the fake data simulation).

Build a CI and Check It

```
b.hat <- coef(lm1)[2]
b.se <- se.coef(lm1)[2]
cover.68 <- abs(beta - b.hat) < b.se # this will be TRUE or FALSE
cover.95 <- abs(beta - b.hat) < 2 * b.se # this will be TRUE or FALSE
cat(paste("68% coverage: ", cover.68, "\n"))

## 68% coverage: FALSE

cat(paste("95% coverage: ", cover.95, "\n"))

## 95% coverage: FALSE
```

It worked once. Need long run coverage.

Long Run Coverage (Normal)

```
n.fake <- 1000
cover.68 <- cover.95 <- rep(NA, n.fake)
for (s in 1:n.fake) {
  y <- alpha + beta * x + rnorm(n, 0, sigma)
  lm.1 <- lm(y ~ x)
  b.hat <- coef(lm.1)[2]
  b.se <- se.coef(lm.1)[2]
  cover.68[s] <- abs(beta - b.hat) < b.se
  cover.95[s] <- abs(beta - b.hat) < 2 * b.se
}
cat(paste("68% coverage: ", mean(cover.68), "\n"))

## 68% coverage: 0.639

cat(paste("95% coverage: ", mean(cover.95), "\n"))

## 95% coverage: 0.88
```

Stat 505 Gelman & Hill, Chapter 8

Fake Data Residuals Sim

```
midterm <- grades[, "Midterm"]
final <- grades[, "Final"]
display.xtable(lm.1 <- lm(final ~ midterm))
```

	Estimate	Std. Error	t value
(Intercept)	64.50	16.98	3.80
midterm	0.70	0.21	3.28

Table: n = 52 rank = 2 resid sd = 14.752 R-Squared = 0.177

```
X <- cbind(1, midterm)
n <- length(final)
```

Nothing simulated yet.

Stat 505 Gelman & Hill, Chapter 8

Long Run Coverage (t_3)

```
for (s in 1:n.fake) {
  y <- alpha + beta * x + rnorm(n, 0, sigma)
  lm.1 <- lm(y ~ x)
  b.hat <- coef(lm.1)[2]
  b.se <- se.coef(lm.1)[2]
  cover.68[s] <- abs(beta - b.hat) < qt(0.84, n - 2) *
    b.se
  cover.95[s] <- abs(beta - b.hat) < qt(0.975, n - 2) *
    b.se
}
cat(paste("68% coverage: ", mean(cover.68), "\n"))

## 68% coverage: 0.652


cat(paste("95% coverage: ", mean(cover.95), "\n"))

## 95% coverage: 0.948
```

Stat 505 Gelman & Hill, Chapter 8

Residual Relationships

```
par(mfrow = c(1, 2))
plot(lm.1, which = 1, add.smooth = F)
plot(final, resid(lm.1), main = "Residuals vs observed")
```




Residuals are orthogonal to Fits, not to Observed.
Make a demo to convince us that's right.

Stat 505 Gelman & Hill, Chapter 8

Residuals Relationships 2

```
a <- 65; b <- 0.7; sigma <- 15; par(mfrow=c(1,2))
y.fake <- a + b*midterm + rnorm(n, 0, sigma)
lm.fake <- lm(y.fake ~ midterm)
plot(lm.fake, which=1, main="Residuals vs Fitted", add.smooth=F, sub="")
abline(h=0, col="grey")
plot(y.fake, resid(lm.fake), main="Residuals vs observed")
abline(h=0, col="grey") ## Pattern seen is not a failure of the model.
```




Stat 505 Gelman & Hill, Chapter 8

§8.3 Simulated Data Compared to Actual Data

Speed of light data (Newcomb 1882) has some outliers.

```
hist((light$time - 24.8) * 1000, breaks = 35, xlab = "Light Travel Time",
     main = "")
```



Fit a mean assuming (poor idea) normality.

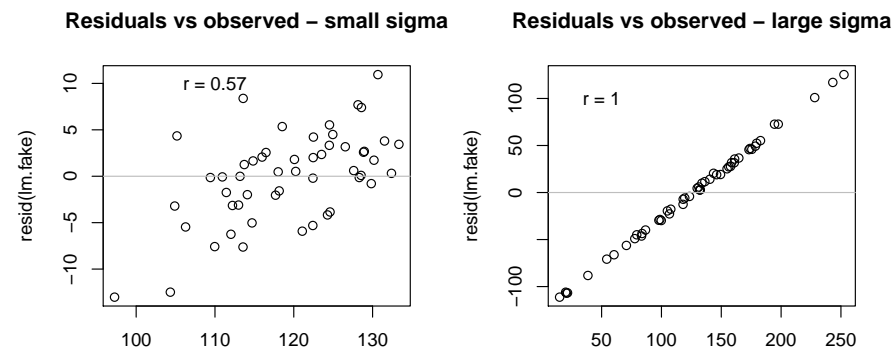
```
display.xtable(light.fit <- lm(I((time - 24.8) * 1000) ~
  1, light))
```

	Estimate	Std. Error	t value
(Intercept)	26.21	1.32	19.82

Stat 505 Gelman & Hill, Chapter 8

Residuals Relationships 3

```
par(mfrow=c(1,2))
y.fake <- a + b*midterm + rnorm(n, 0, 5)
lm.fake <- lm(y.fake ~ midterm)
plot(y.fake, resid(lm.fake), main="Residuals vs observed - small sigma")
abline(h=0, col="grey"); text(110,10, paste("r =", round(cor(y.fake, resid(lm.fake)), 2)))
y.fake <- a + b*midterm + rnorm(n, 0, 50)
lm.fake <- lm(y.fake ~ midterm)
plot(y.fake, resid(lm.fake), main="Residuals vs observed - large sigma")
abline(h=0, col="grey"); text(50,100, paste("r =", round(cor(y.fake, resid(lm.fake)), 2)))
```



Stat 505 Gelman & Hill, Chapter 8

Simulate from Model

Generate 1000 random coefficient estimates using sampling distribution of $(\hat{\beta}^T \hat{\sigma})$

```
sim.light <- sim(light.fit, 1000)
n <- nrow(sim.light)
```


For each simulated coefficient, generate random y_{rep} and plot 50 of them, comparing to the actual data.

```
y.rep <- sapply(1:1000, function(ndx) rnorm(n, sim.light@coef[ndx],
  sim.light@sigma[ndx]))
```

Stat 505 Gelman & Hill, Chapter 8

Plot Sims versus Actual

```
plot(c(-45, 55), c(0, 0.08), type = "n", xlab = "Travel Time",
     main = "")
plot.myDensity <- function(y, ...) {
  dense <- density(y)
  lines(dense$x, dense$y, ...)
}
for (s in 1:50) plot.myDensity(y.rep[, s], col = "grey")
plot.myDensity((light$time - 24.8) * 1000, lwd = 2)
```



Grey lines don't ever have the two bumps, otherwise they have too much spread.

Stat 505 Gelman & Hill, Chapter 8


Conclusion of Newcomb's Data

Density plots and numerical summaries agree:
Newcomb's data don't come from a normal distribution. Could be a contaminated mixture, or a long-tailed, or ...

Numerical Summaries

With simple data, plot shows the observed data are quite different from simulated data. More of a challenge with complex data.
Idea: find a numerical summary $T(y)$ which highlights some aspect of the data. Compute it on each y_{rep} . In this case, the min is a good candidate.

```
Test <- function(y) min(y)
test.rep <- apply(y.rep, 2, Test)
hist(c(min((light$time - 24.8) * 1000), test.rep), main = "Histogram of T",
     xlab = "Minimum travel time")
```



Stat 505 Gelman & Hill, Chapter 8

Count Data Example

Model number of roaches caught as $y_i \sim \text{Poisson}(u_i \exp(X_i \beta))$ where \mathbf{X} has predictors: pre-treatment roach level an indicator for treatment and an indicator for "senior" (building restricted to elderly) and an intercept. Offset u_i is "number of trap days" and $\log(u_i)$ acts like another predictor with coefficient set to be one.

```
roach.glm1 <- glm(y ~ roach1 + treatment + senior, roaches,
                  family = poisson, offset = log(exposure2))
display.xtable(roach.glm1)
```

	Estimate	Std. Error	z value
(Intercept)	3.09	0.02	145.49
roach1	0.01	0.00	78.69
treatment	-0.52	0.02	-20.89
senior	-0.38	0.03	-11.37

Table: n = 262 rank = 4 Resid Deviance = 11429.467

Stat 505 Gelman & Hill, Chapter 8

Stat 505 Gelman & Hill, Chapter 8

Count Data Example: Random Poisson

Generate one set of random Poissons from this model.

```
X <- model.matrix(roach.glm1)
y.hat1 <- roaches$exposure2 * exp(X %*% coef(roach.glm1))
y.rep1 <- rpois(nrow(roaches), y.hat1)
xtable(matrix(table(roaches$y)[1:10], 1, 10, dimnames = list(
  0:9)))
```

	0	1	2	3	4	5	6	7	8	9
1	94	20	11	10	7	7	3	6	3	2

```
xtable(matrix(table(y.rep1)[1:10], 1, 10, dimnames = list(
  names(table(y.rep1))[1:10])))
```

	0	4	5	6	7	8	9	10	11	12
1	1	1	2	6	4	10	14	18	10	13

Data has way more zeroes than random Poisson.

Stat 505

Gelman & Hill, Chapter 8

Overdispersed Poisson

```
roach.glm2 <- update(roach.glm1, family = quasipoisson)
display.xtable(roach.glm2)
```

	Estimate	Std. Error	t value
(Intercept)	3.09	0.17	17.98
roach1	0.01	0.00	9.73
treatment	-0.52	0.20	-2.58
senior	-0.38	0.27	-1.41

Table: n = 262 rank = 4 Resid Deviance = 11429.467

Stat 505

Gelman & Hill, Chapter 8

Count Data Example Simulation

Generate 1000 random Poissons from this model.

```
roach.sim1 <- sim(roach.glm1, 1000)
y.rep <- sapply(1:1000, function(ndx) rpois(nrow(roaches),
  roaches$exposure2 * exp(X %*% roach.sim1@coef[ndx, ])))
Test <- function(y) mean(y == 0)
Test(roaches$y)

## [1] 0.359

summary(test.rep <- apply(y.rep, 2, Test))

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.00064 0.00000 0.00763
```

How unusual is the real data? Compute a p-value.

```
mean(test.rep >= Test(roaches$y))

## [1] 0
```

Too few zeroes. One fix: add in overdispersion.

Stat 505

Gelman & Hill, Chapter 8

Overdispersed Poisson Simulated Data

```
roach.sim2 <- sim(roach.glm2, 1000)
y.hat <- sapply(1:1000, function(ndx) roaches$exposure2 *
  exp(X %*% roach.sim2@coef[ndx, ]))
y.rep <- sapply(1:1000, function(ndx) rnegbin(nrow(roaches),
  y.hat[, ndx], y.hat[, ndx]/(65.4 - 1)))
summary(test.rep <- apply(y.rep, 2, Test))

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.168 0.290 0.317 0.319 0.351 0.469
```

Is 35% zeroes unusual?


Stat 505

Gelman & Hill, Chapter 8

Autocorrelated Data

To use OLS regression, errors must be independent. With time series data, it's common to find that each point is correlated with its neighboring points. Positive AR: highs follow highs, lows follow lows.

```
year <- unemploy$year
y <- unemploy$unemployed.pct
plot(year, y, type = "l", ylab = "unemployment", xlab = "year",
     yaxs = "i", ylim = c(0, max(y) * 1.05), yaxt = "n",
     mgp = c(2, 0.5, 0), cex.axis = 1.2, cex.lab = 1.2)
axis(2, c(0, 5, 10), paste(c(0, 5, 10), "%", sep = ""),
     mgp = c(2, 0.5, 0), cex.axis = 1.2)
```



Stat 505

Gelman & Hill, Chapter 8

Simulate Autocorrelated Data

```
b.hat <- coef(lm.lag) # vector of 2 regression coefs
s.hat <- sigma.hat(lm.lag) # residual sd
n.sims <- 1000
y.rep <- array(NA, c(n.sims, n <- length(y)))
for (s in 1:n.sims) {
  y.rep[s, 1] <- y[1]
  for (t in 2:n) {
    prediction <- c(1, y.rep[s, t - 1]) %*% b.hat
    y.rep[s, t] <- rnorm(1, prediction, s.hat)
  }
}
```

Stat 505

Gelman & Hill, Chapter 8

Autocorrelated Data Fit

Create "lagged y" to predict current year from previous year.

```
y.lag <- c(NA, y)[1:58]
lm.lag <- lm(y ~ y.lag)
display.xtable(lm.lag)
```

	Estimate	Std. Error	t value
(Intercept)	1.43	0.50	2.84
y.lag	0.75	0.09	8.61

Table: n = 57 rank = 2 resid sd = 0.986 R-Squared = 0.574

Is this a good fit? Simulate to see.

Stat 505

Gelman & Hill, Chapter 8

Simulate Autocorrelated Data 2

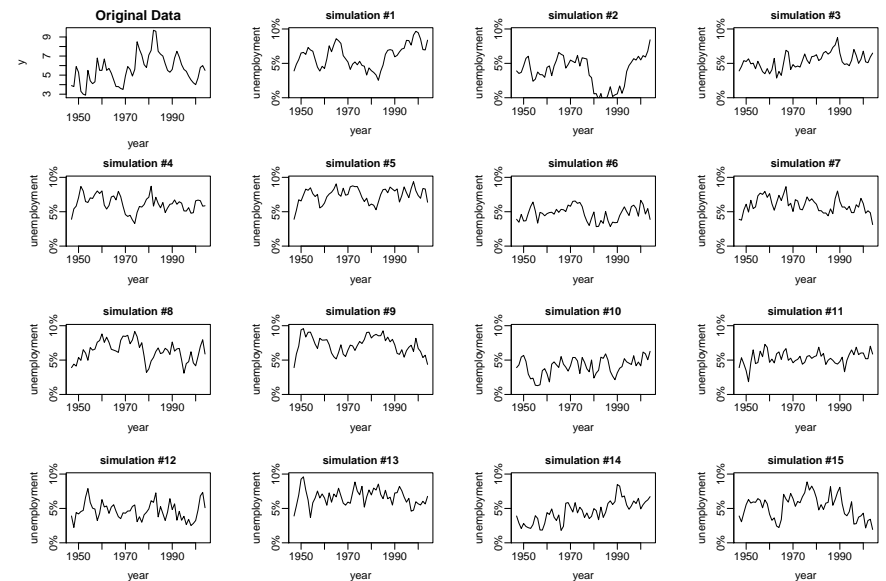
```
## Including uncertainty in the estimated parameters
lm.lag.sim <- sim(lm.lag, n.sims) # simulations of beta and sigma
for (s in 1:n.sims) {
  y.rep[s, 1] <- y[1]
  for (t in 2:n) {
    prediction <- c(1, y.rep[s, t - 1]) %*% lm.lag.sim@coef[s,
    ]
    y.rep[s, t] <- rnorm(1, prediction, lm.lag.sim@sigma[s])
  }
}
```

Stat 505

Gelman & Hill, Chapter 8

Plot of simulated unemployment rate series

```
par(mfrow = c(4, 4), mar = c(4, 4, 2, 2))
plot(year, y, type = "l", main = "Original Data")
for (s in 1:15) {
  plot(year, y.rep[s, ], type = "l", ylab = "unemployment",
        xlab = "year", yaxs = "i", ylim = c(0, max(y) *
        1.05), yaxt = "n", mgp = c(2, 0.5, 0), main = paste("simulation",
        s, sep = ""), cex.main = 0.95)
  axis(2, c(0, 5, 10), paste(c(0, 5, 10), "%", sep = ""),
        mgp = c(2, 0.5, 0))
}
```



Simulate Autocorrelated Data 4

```
Test <- function(y) {
  ## count # of switches in sign of difference
  n <- length(y)
  y.lag <- c(NA, y)[1:58]
  y.lag2 <- c(NA, NA, y)[1:58]
  sum(sign(y - y.lag) != sign(y.lag - y.lag2), na.rm = TRUE)
}
## Test (y) ## 23 switches in original data
n.sims <- 1000
for (s in 1:n.sims) {
  test.rep[s] <- Test(y.rep[s, ])
}
mean(test.rep > Test(y))

## [1] 0.972

quantile(test.rep, c(0.025, 0.05, 0.5, 0.95, 0.975))

## 2.5% 5% 50% 95% 97.5%
## 23 24 31 36 38
```